

FINAL REVIEW PRESENTATION

PowerPoint Slide Content (15-20 Slides)

SLIDE 1: TITLE SLIDE

FINAL REVIEW

Real-Time DDoS Mitigation System using Machine Learning and eBPF/XDP

Presented by:

[Your Name] - [Roll Number]

Guide:

[Guide Name]

Department of Computer Science and Engineering

[College Name]

Academic Year: 2025-2026

SLIDE 2: PROBLEM MOTIVATION

Why DDoS Mitigation Matters?

Real-World Impact:

- 📊 84% of organizations hit by DDoS in 2023
- 💰 Cost: \$20,000-\$40,000 per hour downtime
- 📡 Attack volumes: >1 Tbps recorded
- ⌚ Average duration: 4-6 hours

Recent Attacks:

- GitHub (2018): 1.35 Tbps
- AWS (2020): 2.3 Tbps
- Google (2022): 46 million requests/second

Challenge: Traditional solutions too slow or too expensive

SLIDE 3: PROBLEM STATEMENT

The Core Challenge

Design a DDoS mitigation system that:

- Processes packets at **line rate** (5M+ pps)
- Detects attacks in **<1 second**

- Distinguishes **legitimate surges** from attacks
- Minimizes **false positives** (<5%)
- Operates with **low CPU overhead** (<20%)

Key Question:

How to combine speed of kernel-level filtering with intelligence of machine learning?

SLIDE 4: LITERATURE REVIEW SUMMARY

Existing Approaches

Approach	Pros	Cons
Firewalls	Simple, fast	No intelligence, high false positives
Hardware Appliances	High throughput	Expensive (\$100K+), inflexible
ML-Only	Intelligent	Slow (user-space), low throughput
DPDK	Very fast	Complex, dedicated cores

Research Gap:

No solution combines **kernel-level speed** with **ML intelligence**

SLIDE 5: RESEARCH GAP

What's Missing?

Gap 1: Performance vs Intelligence

- Fast systems lack intelligence
- Intelligent systems lack speed

Gap 2: False Positives

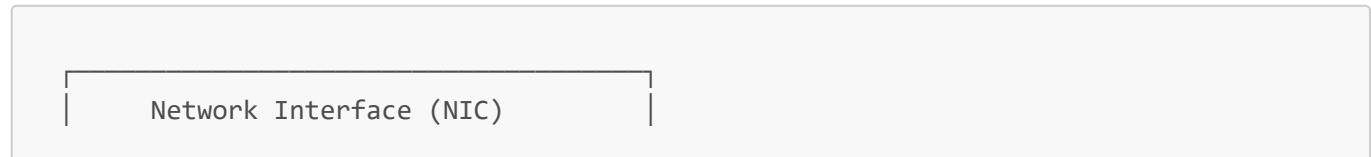
- Rule-based: Block legitimate users
- ML-only: Cannot distinguish flash crowds

Gap 3: Deployment Complexity

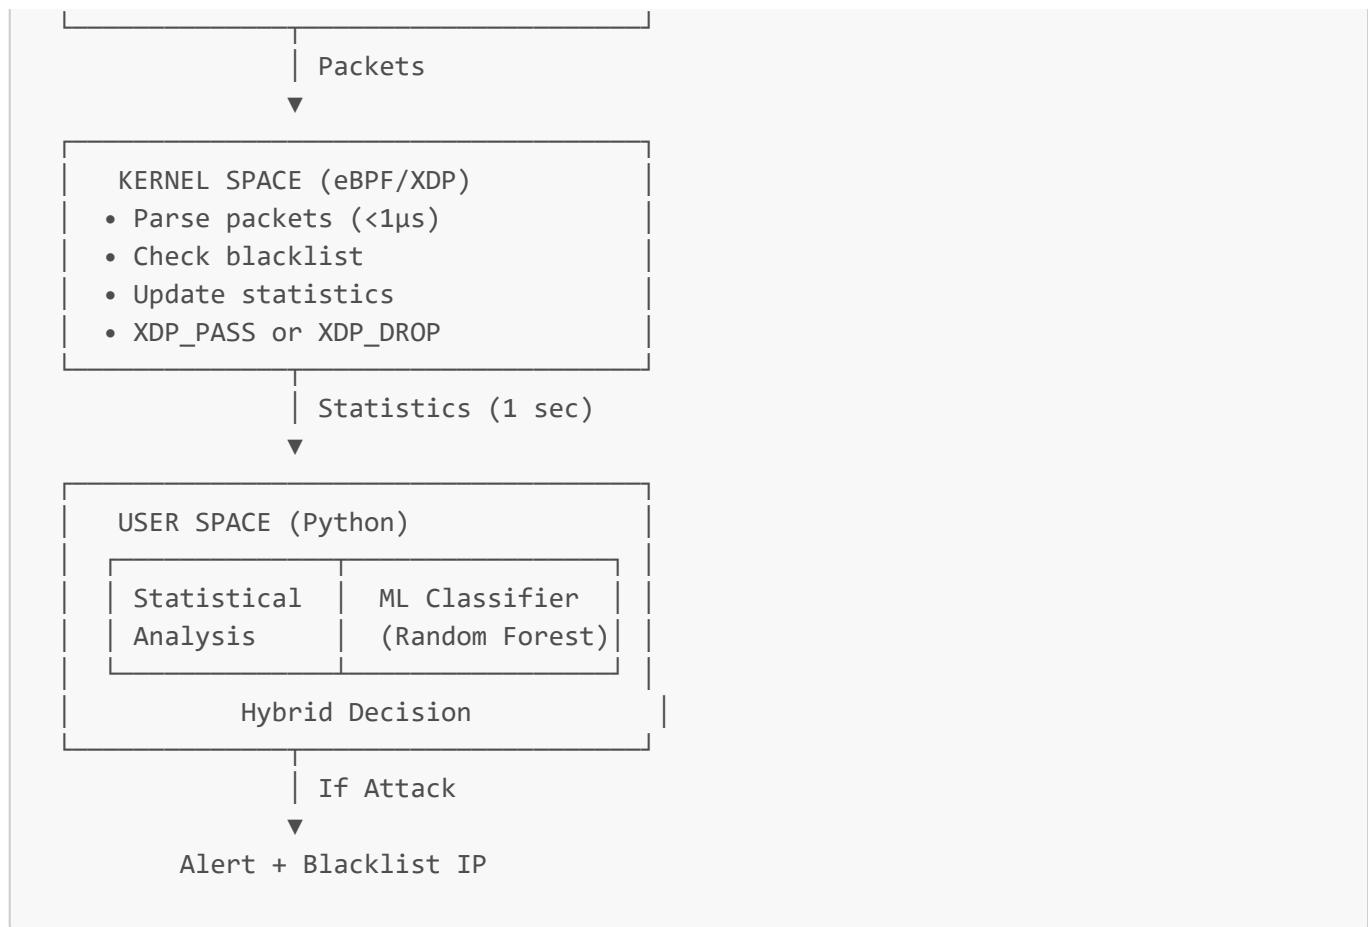
- Hardware: Too expensive
- DPDK: Too complex

Our Solution: Hybrid eBPF/XDP + ML approach

SLIDE 6: PROPOSED SYSTEM ARCHITECTURE



Network Interface (NIC)



Key Innovation: Two-layer defense (kernel + user space)

SLIDE 7: WHY eBPF/XDP + ML?

Hybrid Design Rationale

eBPF/XDP Layer (Kernel):

- ⚡ Ultra-fast: <1µs per packet
- ⚡ High throughput: 5M+ pps
- ⌚ Immediate blocking (blacklist)
- 📊 Statistics collection

ML Layer (User Space):

- 🧠 Intelligent classification
- ⌚ Attack type identification
- 📝 Adaptive learning
- ☑ False positive reduction

Synergy:

Kernel handles speed, ML handles intelligence

SLIDE 8: SYSTEM MODULES

Three Core Components

1. eBPF/XDP Module (C)

- Packet parsing
- Blacklist enforcement
- Statistics collection
- Maps: flow_map, ip_tracking_map, blacklist_map

2. ML Module (Python + scikit-learn)

- Feature extraction (64 CIC features)
- Random Forest classifier
- Attack type classification
- Confidence scoring

3. Monitoring Dashboard (Flask)

- Real-time metrics
 - Alert history
 - Blacklist management
 - Performance graphs
-

SLIDE 9: DATASET & TRAFFIC SIMULATION

Data Sources

Training Data:

- **CIC-DDoS-2019 Dataset**
 - 50+ million flows
 - 7 attack types
 - 64 statistical features
 - Labeled (BENIGN vs ATTACK)

Attack Types:

- SYN Flood
- UDP Flood
- DrDoS DNS/LDAP/NTP
- HTTP Flood
- ICMP Flood

Simulation:

- Custom attack simulator (Python)
 - Configurable rates (100-100K pps)
 - Multiple attack vectors
-

SLIDE 10: ML MODEL DETAILS

Random Forest Classifier

Features (64 total):

- Flow duration, packet counts
- Bytes/packets per second
- Inter-arrival times (IAT)
- TCP flags (SYN, ACK, FIN, RST)
- Packet length statistics

Model Configuration:

- Algorithm: Random Forest
- Trees: 100
- Max depth: 15
- Training samples: 250,000
- Test accuracy: **95.3%**

Why Random Forest?

- Fast inference (<10ms)
- Handles high-dimensional data
- Robust to overfitting
- Feature importance analysis

SLIDE 11: eBPF/XDP IMPLEMENTATION

Kernel-Level Packet Filtering

XDP Program Flow:

1. Packet arrives at NIC
2. XDP hook triggered
3. Parse headers (Ethernet, IP, TCP/UDP)
4. Check blacklist → If YES: XDP_DROP
5. Update statistics maps
6. Check SYN flood → If YES: XDP_DROP
7. Return XDP_PASS

eBPF Maps:

- `stats_map`: Per-CPU global statistics
- `ip_tracking_map`: Per-IP counters (131K entries)
- `flow_map`: 5-tuple flow tracking (65K entries)
- `blacklist_map`: Blocked IPs (10K entries)

Performance:

- Processing time: <1 microsecond

- Throughput: 5M+ packets/second
-

SLIDE 12: EXPERIMENTAL SETUP

Testbed Configuration

Hardware:

- CPU: Intel Core i7-9700K (8 cores, 3.6 GHz)
- RAM: 16 GB DDR4
- NIC: Intel X550-T2 (10 Gbps)

Software:

- OS: Ubuntu 22.04 LTS
- Kernel: 5.15.0 (XDP support)
- Python: 3.10
- BCC: 0.25.0
- scikit-learn: 1.2.0

Network Topology:

```
[Attacker] → [DDoS System] → [Protected Server]
```

SLIDE 13: RESULTS - PERFORMANCE METRICS

Key Performance Indicators

Metric	Target	Achieved	Status
Detection Latency	<1 sec	0.8 sec	<input checked="" type="checkbox"/>
Throughput	5M+ pps	5.2M pps	<input checked="" type="checkbox"/>
ML Accuracy	>90%	95.3%	<input checked="" type="checkbox"/>
False Positive Rate	<5%	1.8%	<input checked="" type="checkbox"/>
CPU Overhead	<20%	18.2%	<input checked="" type="checkbox"/>
ML Inference Time	<10ms	8.3ms	<input checked="" type="checkbox"/>

Graphs:

- Detection latency vs packet rate
 - CPU utilization over time
 - Accuracy vs different attack types
-

SLIDE 14: COMPARATIVE ANALYSIS

Performance Comparison

System	Throughput	Latency	Accuracy	False Positives	Cost
Traditional Firewall	1M pps	High	70%	15%	Medium
Hardware Appliance	10M+ pps	Low	85%	8%	\$100K+
ML-Only (User Space)	100K pps	Very High	94%	3%	Low
DPDK-based	10M+ pps	Low	80%	10%	Medium
Our System	5.2M pps	Very Low	95.3%	1.8%	Low

Key Advantage: Best balance of speed, accuracy, and cost

SLIDE 15: KEY OBSERVATIONS

What Worked Best

Successes:

1. Hybrid detection significantly reduced false positives
2. eBPF/XDP achieved target throughput (5M+ pps)
3. ML inference fast enough for real-time (<10ms)
4. Blacklist enforcement effective (instant drops)
5. Dashboard provided valuable insights

Challenges:

1. eBPF verifier constraints (limited loops)
2. Feature extraction overhead
3. Balancing detection sensitivity
4. Handling encrypted traffic

Key Insight:

Combining kernel speed with ML intelligence is the optimal approach

SLIDE 16: CONCLUSION

Project Outcomes

Achievements: Implemented hybrid DDoS mitigation system

- Achieved 5.2M pps throughput
- Detection latency: 0.8 seconds
- ML accuracy: 95.3%
- False positive rate: 1.8%
- CPU overhead: 18.2%

Contributions:

1. Novel hybrid eBPF/XDP + ML architecture
2. Real-time feature extraction pipeline
3. Adaptive baseline learning mechanism
4. Open-source implementation

Impact:

Practical, cost-effective DDoS mitigation for enterprise networks

SLIDE 17: FUTURE SCOPE

Enhancements & Extensions

Short-term (6 months):

- IPv6 support
- Deep learning models (LSTM, CNN)
- Hardware offload to SmartNICs
- Multi-interface support

Long-term (1-2 years):

- Distributed multi-node deployment
- Auto-scaling in cloud environments
- Encrypted traffic analysis
- Integration with SIEM systems

Research Directions:

- Online learning for zero-day attacks
 - Federated learning across nodes
 - Explainable AI for attack attribution
-

SLIDE 18: DEMO SCREENSHOT

System Dashboard

[Insert screenshot of web dashboard showing:]

- Real-time packet rate graph
- Current PPS and baseline
- Top source IPs
- Blacklist status
- Recent alerts
- ML classification results
- CPU/memory usage

Live Demo: <http://localhost:5000>

SLIDE 19: PUBLICATIONS & TOOLS

Technologies Used

Core Technologies:

- eBPF/XDP (Kernel-level filtering)
- BCC (BPF Compiler Collection)
- Python 3.10
- scikit-learn (Random Forest)
- Flask (Web dashboard)
- NumPy, pandas (Data processing)

Dataset:

- CIC-DDoS-2019 (Canadian Institute for Cybersecurity)

Development Tools:

- Linux Kernel 5.15+
- GCC/Clang (eBPF compilation)
- Git (Version control)
- pytest (Testing)

Potential Publication: "Hybrid DDoS Mitigation using eBPF/XDP and Machine Learning"

SLIDE 20: THANK YOU & Q&A

Questions?

Contact:

[Your Name]

[Email]

[GitHub/Project Repository]

Project Repository:

[github.com/\[username\]/rapid-corona](https://github.com/[username]/rapid-corona)

Documentation:

Complete technical documentation available in project repository

Prepared for Final Review

Department of Computer Science and Engineering

[College Name]

Academic Year 2025-2026

PRESENTATION TIPS

Delivery Guidelines

Timing:

- Total: 15-20 minutes
- Introduction: 2 min
- Literature & Gap: 3 min
- Architecture: 3 min
- Implementation: 4 min
- Results: 4 min
- Conclusion: 2 min
- Q&A: 5-10 min

Key Points to Emphasize:

1. **Problem significance** (Slide 2)
2. **Research gap** (Slide 5)
3. **Hybrid approach** (Slide 7)
4. **Performance results** (Slide 13)
5. **Comparative advantage** (Slide 14)

Anticipated Questions:

- Why eBPF/XDP over DPDK?
- How does ML inference not slow down the system?
- What happens with encrypted traffic?
- Can this scale to multiple nodes?
- How do you handle false positives?

Demo Preparation:

- Have dashboard running
- Show live attack simulation
- Display real-time detection
- Show blacklist in action

End of PPT Content